

SEE EXACTLY WHAT YOUR AGENT ACTUALLY DID

RuleReceipt checks a coding agent's session against your own rules — CLAUDE.md or AGENTS.md — and shows precisely what was followed, what wasn't, and why. Runs on your machine. Nothing sent anywhere by default.

```
$ npx rulereceipt check
```

THE PROBLEM, DOCUMENTED

#18660 CLAUDE.md instructions are read but not reliably followed

#15443 Claude ignores explicit instructions while claiming to understand them

#24318 Claude ignores explicit instructions and acts without approval

A developer told Claude explicitly: "Don't deploy."
It ran a deploy command over production anyway, corrupting the database.

WHY IT MATTERS FOR TEAMS

Vendor security reviews and DPAs ask how you actually develop. A security questionnaire asking about secrets handling or input validation is really asking: does your process actually enforce this? Most teams answer from memory, not evidence.

Protects the company from what a session did silently — and protects the dev, too. Evidence, not blame.

WHAT WE DON'T DO

— Nothing leaves your machine

Runs locally, always. Only server call is your own Claude key to Anthropic — nothing else.

— We don't see your data

No server of ours sits in the path.
What you check stays yours.

— We never touch settings.json

Read-only, always. No hooks, no modifications, without you asking.

"I rely on documented rules being followed as my QA process — CLAUDE.md isn't a nice-to-have, it's how I maintain engineering standards without being an engineer."

— a developer who built an 82,000-line app through Claude Code alone

FREE during early access, then \$1 / ~a month of daily checks

rulereceipt.dev →