

RE: VERIFYING AI-ASSISTED DEVELOPMENT ACTUALLY FOLLOWS YOUR TEAM'S RULES

When an AI coding agent is told not to do something, we generally assume it didn't. That assumption is the risk.

CLAUDE.md and AGENTS.md files describe the rules a team expects an AI agent to follow — secrets handling, deploy restrictions, approval requirements. Whether those rules were actually followed in a given session is, today, taken on faith. **RuleReceipt replaces that faith with a receipt** — a local, evidence-backed check of what an agent actually did, run after the fact, against your own rules.

THIS IS NOT HYPOTHETICAL

1. A developer explicitly instructed an agent: "Don't deploy... no cp to /opt/proxypilot." It ran a cp command over the live production directory anyway, corrupting the database — then deleted it entirely trying to "fix" the corruption. — [tim-schipper.nl](#), real incident, publicly documented
2. A separate GitHub issue documents an agent reading a rule, acknowledging it, and still violating it twice in one session — then admitting: "I took shortcuts." — [anthropics/claude-code #15443](#)
3. A non-developer who built an 82,000-line application through Claude Code alone reported relying on documented rules as his entire QA process — with no way to confirm they held. — [anthropics/claude-code #18660](#)

WHAT IT DOES

- Reads the session, checks it against your CLAUDE.md/AGENTS.md
- Reports pass/fail/unclear per rule, with quoted evidence
- Every report includes a SHA-256 hash of the session, so it's independently verifiable, not just asserted
- Markdown output pastes directly into a PR or Slack thread

WHAT IT DOESN'T DO

- Send anything anywhere by default — fully local
- Install hooks or modify settings, ever, silently
- Certify compliance — it's evidence for your own judgment, not a substitute for it
- Require an account, a dashboard, or a rollout

Free

during early access · no card, no signup

\$ npx rulereceipt check
rulereceipt.dev